

Exploratory Data Analysis - Written Exam

André Plancha

andre.plancha@hotmail.com

Computer Science: Applied Data Science

MALMÖ UNIVERSITY

Spring 2025

March 29, 2025

1. Due to large sizes of data sets in practice one needs to do dimension reduction thereafter study the pattern by various clustering techniques. Dimension reduction means loss of certain information whence less accuracy of clustering results in general. In this question, we consider the Iris data set as a four-dimensional data set with the ground truth label: setosa, versicolor, virginica and use the k-means for clustering.

a) First determine three clusters of Iris data set by k-means, and even compute the percentage of correct classified observations.


K-means clustering is an iterative algorithm that aims to partition n observations into k clusters in which each observation belongs to the cluster with the nearest mean. Using scikit-learn's `KMeans`, the following was obtained:

	Cluster 1	Cluster 2	Cluster 3
setosa	0 (0%)	50 (100%)	0 (0%)
versicolor	47 (77%)	0 (0%)	3 (8%)
virginica	14 (23%)	0 (0%)	36 (92%)
Total	61	50	39

From this, we can see that the clusters obtained are able to separate the classes quite well, with 100% of setosa on one cluster, 77% of versicolor on another, and 92% of virginica on the last one. This leads us to a total of 82% of correct classified observations.

b) Now apply first the factor analysis method to reduce the Iris data set to two-dimensional, then determine three clusters of the reduced data set by k-means. What is then the percentage of correct classified observations? Even visualize the reduced data set with original label respectively by the clustering.

Factor analysis uses the correlation structure amongst observed variables to model a smaller number of unobserved, latent variables known as factors. Using scikit-learn's `FactorAnalysis` and scikit-learn's `KMeans`, the following was obtained (the color in the background represents the cluster division):



	Cluster 1	Cluster 2	Cluster 3
setosa	0 (0%)	50 (100%)	0 (0%)
versicolor	15 (38%)	0 (0%)	35 (57%)
virginica	24 (62%)	0 (0%)	26 (43%)
Total	39	50	61

From this, we can clearly see that the algorithm was able to separate setosas from the other classes, but it had a hard time separating versicolors from virginicas. This leads us to a total of $\approx 72.67\%$ of correct classified observations. This is a decrease in the accuracy of the classification, which is expected when reducing the dimensionality of the data.

2. Consider the minimum spanning trees method (MST) and lung cancer data set LungA.

a) First apply MST to find two clusters, then find the percentage of correct clustered genes using the original label of cancer cell: Normal, Cancer (Small-cell lung carcinomas, Non-small cell lung carcinomas) as ground truth.

As I interpreted it, this question asks us to:

1. Transform LungA into a undirected weighted complete graph, where the weight of the edge between two nodes is the euclidean distance between them;

2. Transform that graph into its minimum spanning tree, using for example Kruskal's algorithm;
3. Remove the longest edge from the tree, thus dividing it into two clusters;
4. Compare the clusters obtained with the original labels.

LungA is a dataset that seems to have 203 observations and 3312 unnamed numerical features, and each observation has a label between AD, SQ, COID, NL, and SCLC. Grouping AD, SQ, COID and SCLC as "Cancer", using the described steps above and using SciPy's `minimum_spanning_tree` and `connected_components`, the following was obtained:

	Cluster 1	Cluster 2
Cancer	185 (92%)	1 (100%)
Normal	17 (8%)	0 (0%)
Total	202	1

From this, we can see that the algorithm was not able to separate the normal cells from the cancer cells, since the smaller cluster has only one Cancer observation. Because of the nature of this algorithm, we can conclude that this observation might be an outlier, showing the method's robustness in outlier detection. Assuming cluster 1 is the cancer cluster, the percentage of correct classified observations is 91.13%.

b) Apply MST once again to find three clusters, then compare the correct clustered genes using the original label: Nonsmall cell lung carcinomas, Normal, Small-cell lung carcinomas as ground truth.

This time, AD, SQ and COID were grouped as "Nonsmall cell lung carcinomas", SCLC as "Small-cell lung carcinomas", and NL as "Normal". Using the same steps as before, the following was obtained:

	Cluster 1	Cluster 2	Cluster 3
Nonsmall cell lung carcinomas	178 (89%)	1 (100%)	1 (100%)
Normal	17 (8%)	0 (0%)	0 (0%)
Small-cell lung carcinomas	6 (3%)	0 (0%)	0 (0%)
Total	201	1	1

The same conclusion can be drawn from this result, where the algorithm was not able to separate the observations based on the labels given, reinforcing the idea that the algorithm is robust to outliers. Additionally, it's important to note that, since the method is a divisive clustering method, cluster 3 is a subset of the previous cluster 1, and increasing the number of clusters will result in this cluster 1 being divided into smaller clusters.

c) Finally, test to find five clusters comparing the correct clustered genes using the label: AD, COID, SQ, NL, SCLC as ground truth. Even comment on these clustering results.

This time, AD, COID, SQ, NL and SCLC were used as labels. Using the same steps as before, the following was obtained:

	Cluster 1	Cluster 2	Cluster 3	Cluster 4
AD	138 (70%)	1 (100%)	0 (0%)	0 (0%)
COID	16 (8%)	0 (0%)	3 (100%)	1 (100%)
NL	17 (9%)	0 (0%)	0 (0%)	0 (0%)
SMCL	6 (3%)	0 (0%)	0 (0%)	0 (0%)
SQ	21 (11%)	0 (0%)	0 (0%)	0 (0%)
Total	198	1	3	1


As in a) and b), the algorithm was not able to separate the observations based on the labels given. This might also be because there's a big imbalance between the number of AD observations, which make up a grand majority of the dataset, making the division difficult to obtain from the data.

3. Model based clustering in chapter 6 (section 6.5) is based on geometrical properties of clusters, e.g. balls and ellipsoids in 3-dim. In this question, you are asked to randomly generate data sets to check the capability of MBC: two balls of different sizes, two ellipsoids of different size with symmetry axis parallel with coordinates axis, three ellipsoids also different sizes with arbitrary symmetry axis.

a) Randomly generate those sub data sets so that they are disjoint (far away from each other), then apply model based clustering to find clusters even compare the estimated clusters with original labels.

For this question, 1000 points of every shape were generated, with the following properties:

	x_c	y_c	z_c	x_r	y_r	z_r	α	β	γ
Ball 1	10	10	10	2	2	2	0°	0°	0°
Ball 2	-10	-10	-10	1	1	1	0°	0°	0°
Ellipsoid 1	0	0	0	3	2	1	0°	0°	0°
Ellipsoid 2	10	0	0	1	1.5	1	0°	0°	0°
Ellipsoid 3	0	10	0	4	5	1	30°	45°	30°
Ellipsoid 4	0	0	10	2	2	1	60°	0°	0°
Ellipsoid 5	0	10	10	3	1	1	60°	20°	45°



Model based clustering is a clustering method based on a finite mixture probability model, using the EM algorithm

b) Randomly generate those sub data sets so that they overlap partly, then apply model based clustering to find clusters even compare the associated clusters with original labels.

4. Find the datasets about inflation, unemployment in member countries of European union during the last ten years (2015 - 2024). Your dataset should contain at least 20 countries. Denote the inflation dataset by EUin and the unemployment dataset by EUun and $EU = [EUin, EUun]$.

Hint You may check the database at the websites: <https://european-union.europa.eu> or <https://ec.europa.eu/eurostat/data/database> or

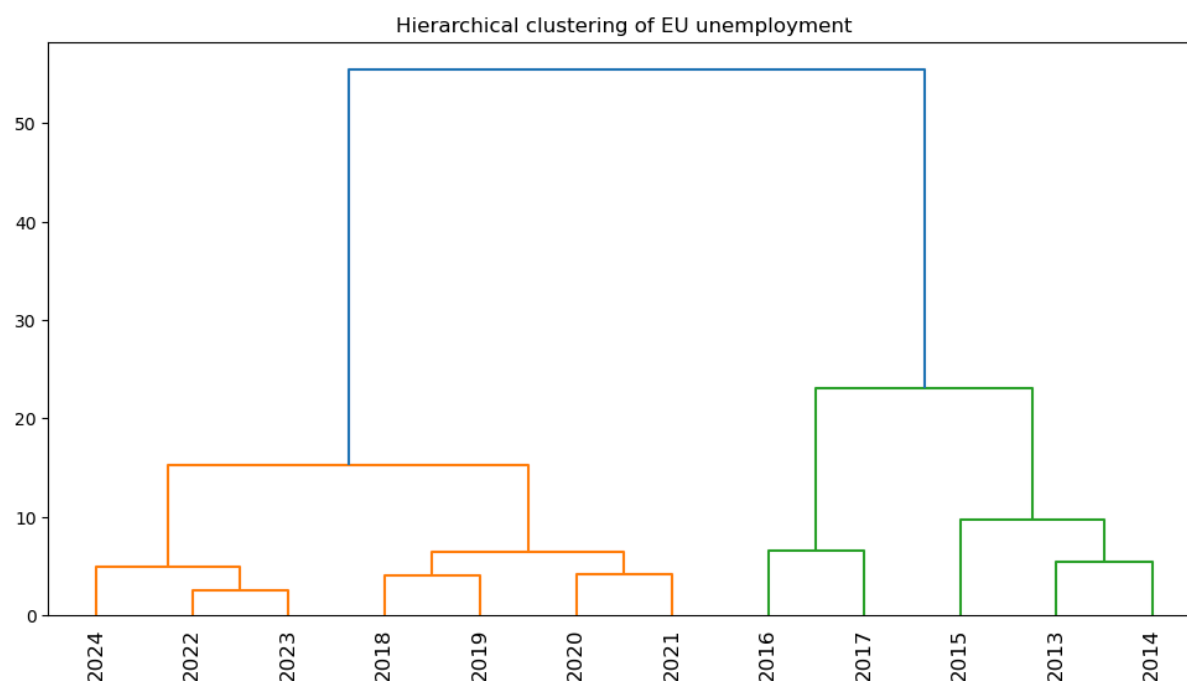
a) Study the unemployment pattern by looking at clusters of EUun using one proper hierarchical methods and evaluate the result by Silhouette plot.

EU_{un} collected is a 12x26 (12 years, 26 EU countries) table with the following structure:

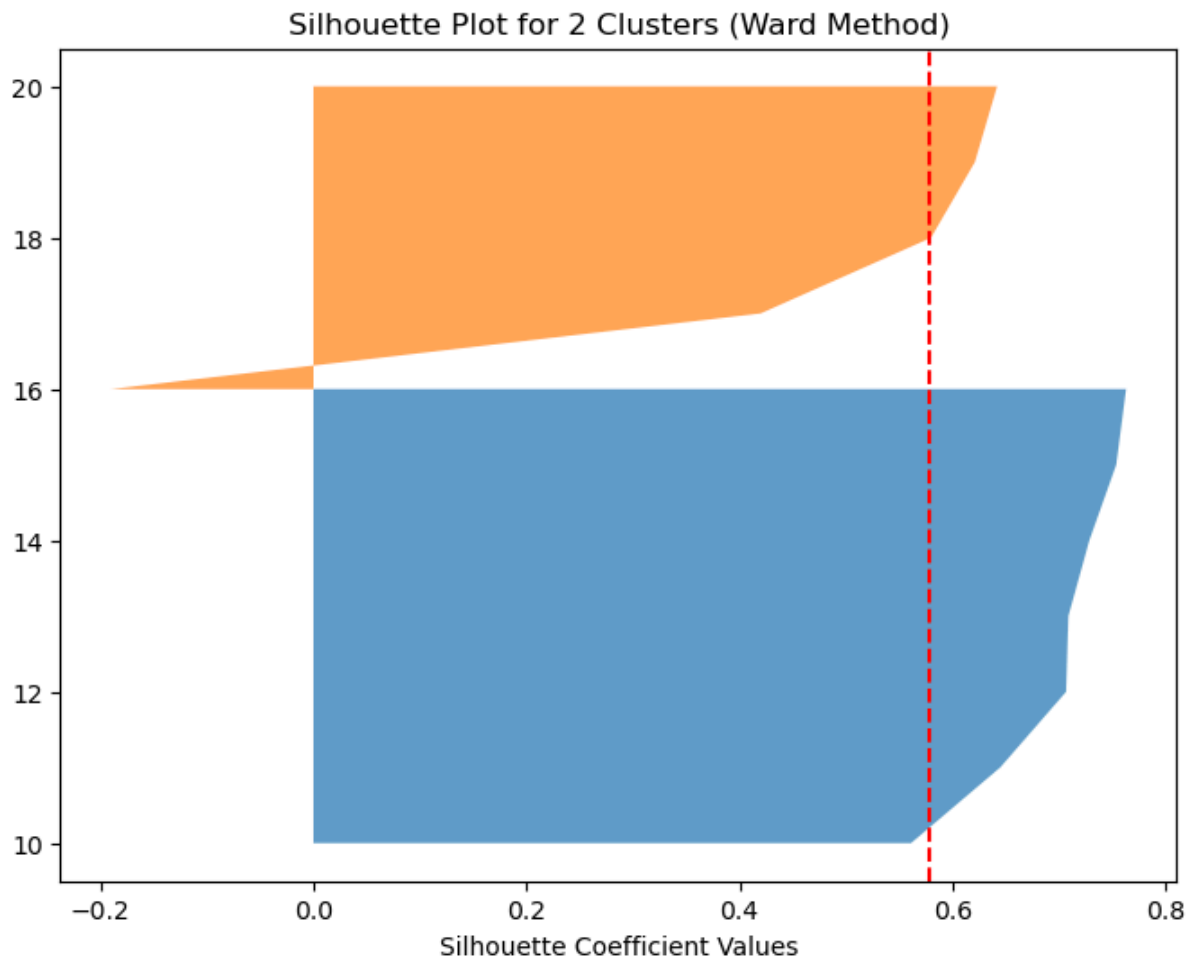
	Belgium	Bulgaria	...	Sweden
2013	8.6	13.9	...	8.1
2014	8.7	12.4	...	8.0
2015	8.7	10.1	...	7.5
...
2024	5.7	4.2	...	8.4

The full dataset is available on github.com/notPlancha/eda-exam/tree/main/data

Hierarchical clustering is a method of clustering analysis that builds a hierarchy of clusters, represented as a tree-like structure called a dendrogram. Using SciPy's linkage with Ward's method, the following dendrogram was obtained:



We can see that the method found a clear divide between years ≤ 2017 and > 2017 in the data, based on the distance between this divide. By defining these 2 as clusters, the following silhouette was produced:

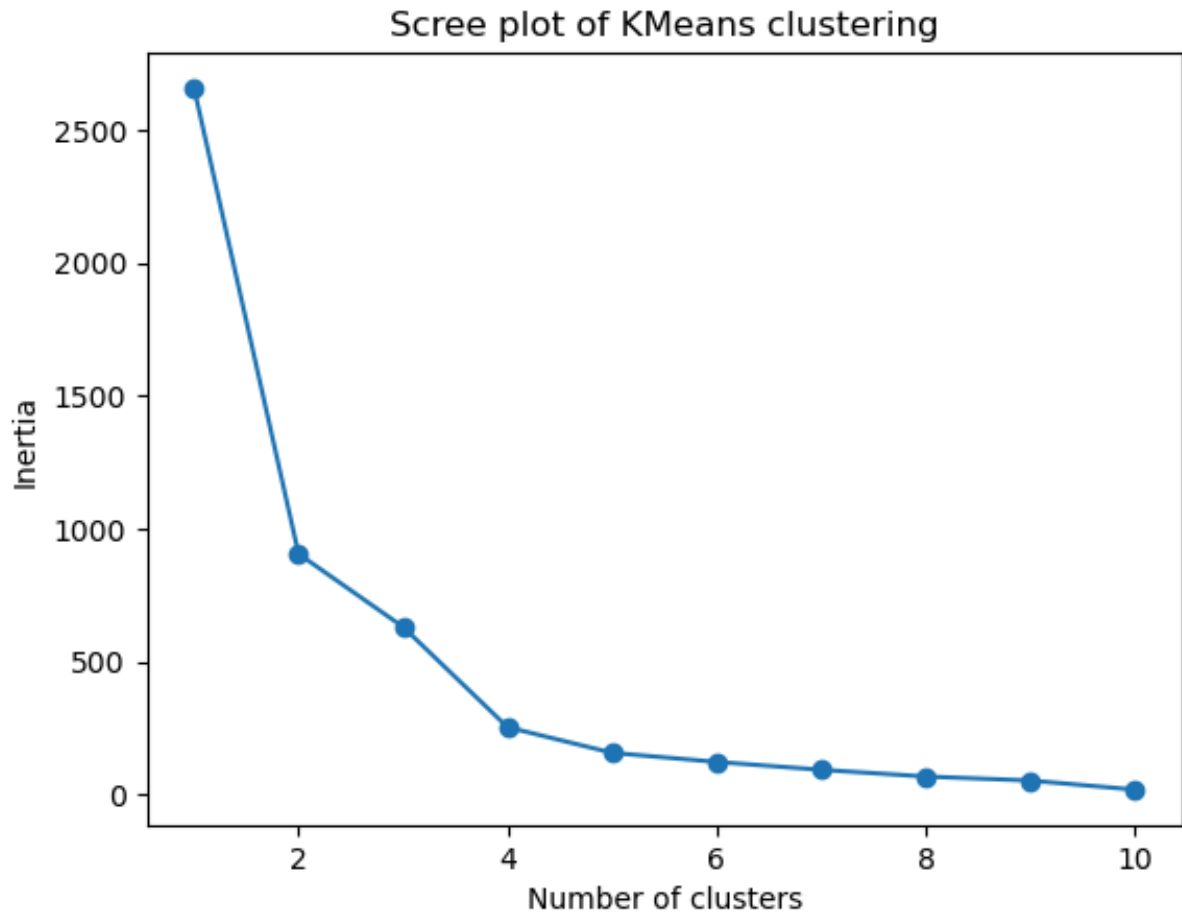


Looking at the silhouette plot, we can see that the clusters are well defined, with a silhouette score close to 0.6. There seems to be a single year that perhaps could be better in the other cluster, but overall the clusters are well defined.

b) Study the joint pattern of unemployment and inflation by looking at clusters of EU using k-means method and evaluate the result by using Dunn index.

EU collected is a 12x52 (12 years, 26*2 EU countries) table. EU_{in} is also available on the same link as before.

To determine the number of clusters that should be used with the k-means algorithm, a scree plot was produced:



By the plot and the elbow method, it seems that 4 clusters should be used. Using scikit-learn's KMeans with 4 clusters, the following Dunn index was obtained:

$$DI_m \approx 0.68$$


The Dunn index is a measure of cluster separation and cluster compactness representing the ratio between the minimum inter-cluster distance and maximum intra-cluster distance: the larger the index, the better the clustering. Although the metric is better as a comparison between cluster results, we can say that since the value is smaller than 1, the clustering is not optimal, but still good.

5. Consider the dataset

$$X = \begin{pmatrix} 3 & 1 & 0.1 \\ 1 & 4 & 0.02 \\ 1 & 1 & 0 \\ 4 & 1 & 0.1 \\ 1.5 & 3 & 0 \\ 0.12 & 2 & 1.9 \\ 0 & 6 & 3.5 \\ 0.03 & 0.5 & 1 \\ 0.1 & 4 & 3 \end{pmatrix}$$


a) Visualize the dataset X via scatter3.

Using plotly's `scatter_3d`, the following was obtained:



b) It's obvious that X contains two clusters: one is on xy-plane and another one is on the yz-plane. Please check if hierarchical Ward's method, spectral method are able to recover these two clusters.


Ward's method is a method of hierarchical clustering that minimizes the within-cluster variance. Using SciPy's `linkage` with Ward's method, and request 2 clusters, the following was obtained:



We can see from the plot that the method was not able to recover these 2 clusters, with 75% of correct classified observations (assuming the xy-plane is cluster 2). This is expected,

since, visually, the variance in the yz-plane seems to be higher than the variance in the xy-plane, and some points between the 2 are very close by.

Spectral clustering is a method that uses the eigenvalues of a similarity matrix to reduce the dimensionality of the data before clustering in a lower dimensional space. Using scikit-learn's `SpectralClustering` with 2 clusters, the following was obtained:



Unexpectedly, the spectral method was also not able to recover the clusters, with the same precision as Ward's method (assuming the yz-plane is cluster 0).